

Retrieval-Augmented Generation, commonly called RAG, combines information retrieval with text generation.

A RAG system first retrieves relevant documents from a knowledge base. The retrieved documents are then provided as context to a language model.

The language model uses the retrieved context to generate an answer. RAG can help ground responses in external information and reduce unsupported answers.

RAG is commonly used for document question answering and knowledge-based AI applications.